

# HR EMPLOYEE ATTRITION ANALYSIS

```
In [1]: # import libraries
import pandas as pd
import numpy as np
```

```
In [2]: df = pd.read_csv('HR-Employee-Attrition.csv')
```

```
In [3]: df.head()
```

Out[3]:

	Age	Attrition	BusinessTravel	DailyRate	Department	DistanceFromHome	Education	EducationField	Emplc
0	41	Yes	Travel_Rarely	1102	Sales	1	2	Life Sciences	
1	49	No	Travel_Frequently	279	Research & Development	8	1	Life Sciences	
2	37	Yes	Travel_Rarely	1373	Research & Development	2	2	Other	
3	33	No	Travel_Frequently	1392	Research & Development	3	4	Life Sciences	
4	27	No	Travel_Rarely	591	Research & Development	2	1	Medical	

5 rows × 35 columns

```
In [4]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1470 entries, 0 to 1469
Data columns (total 35 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Age                                   1470 non-null   int64
1   Attrition                            1470 non-null   object
2   BusinessTravel                       1470 non-null   object
3   DailyRate                            1470 non-null   int64
4   Department                           1470 non-null   object
5   DistanceFromHome                     1470 non-null   int64
6   Education                             1470 non-null   int64
7   EducationField                       1470 non-null   object
8   EmployeeCount                        1470 non-null   int64
9   EmployeeNumber                       1470 non-null   int64
10  EnvironmentSatisfaction               1470 non-null   int64
11  Gender                               1470 non-null   object
12  HourlyRate                           1470 non-null   int64
13  JobInvolvement                       1470 non-null   int64
14  JobLevel                             1470 non-null   int64
15  JobRole                              1470 non-null   object
16  JobSatisfaction                       1470 non-null   int64
17  MaritalStatus                        1470 non-null   object
18  MonthlyIncome                        1470 non-null   int64
19  MonthlyRate                           1470 non-null   int64
20  NumCompaniesWorked                   1470 non-null   int64
21  Over18                               1470 non-null   object
22  OverTime                             1470 non-null   object
23  PercentSalaryHike                    1470 non-null   int64
24  PerformanceRating                    1470 non-null   int64
25  RelationshipSatisfaction              1470 non-null   int64
26  StandardHours                        1470 non-null   int64
```

```
27 StockOptionLevel      1470 non-null   int64
28 TotalWorkingYears      1470 non-null   int64
29 TrainingTimesLastYear  1470 non-null   int64
30 WorkLifeBalance        1470 non-null   int64
31 YearsAtCompany         1470 non-null   int64
32 YearsInCurrentRole     1470 non-null   int64
33 YearsSinceLastPromotion 1470 non-null   int64
34 YearsWithCurrManager   1470 non-null   int64
dtypes: int64(26), object(9)
memory usage: 402.1+ KB
```

```
In [5]: # checking for duplicate values
df.duplicated().sum()
```

```
Out[5]: 0
```

```
In [6]: # summary statistics
df.describe()
```

```
Out[6]:
```

	Age	DailyRate	DistanceFromHome	Education	EmployeeCount	EmployeeNumber	Environment
count	1470.000000	1470.000000	1470.000000	1470.000000	1470.0	1470.000000	
mean	36.923810	802.485714	9.192517	2.912925	1.0	1024.865306	
std	9.135373	403.509100	8.106864	1.024165	0.0	602.024335	
min	18.000000	102.000000	1.000000	1.000000	1.0	1.000000	
25%	30.000000	465.000000	2.000000	2.000000	1.0	491.250000	
50%	36.000000	802.000000	7.000000	3.000000	1.0	1020.500000	
75%	43.000000	1157.000000	14.000000	4.000000	1.0	1555.750000	
max	60.000000	1499.000000	29.000000	5.000000	1.0	2068.000000	

8 rows × 26 columns

```
In [7]: # Check for unique values in the categorical columns

for col in df.columns:
    if df[col].dtype == 'object':
        print(f'{col} unique values are :',df[col].unique())
        print('\n')
```

```
Attrition unique values are : ['Yes' 'No']
```

```
BusinessTravel unique values are : ['Travel_Rarely' 'Travel_Frequently' 'Non-Travel']
```

```
Department unique values are : ['Sales' 'Research & Development' 'Human Resources']
```

```
EducationField unique values are : ['Life Sciences' 'Other' 'Medical' 'Marketing' 'Technical Degree'
'Human Resources']
```

```
Gender unique values are : ['Female' 'Male']
```

```
JobRole unique values are : ['Sales Executive' 'Research Scientist' 'Laboratory Technician'
'Manufacturing Director' 'Healthcare Representative' 'Manager']
```

```
'Sales Representative' 'Research Director' 'Human Resources']
```

```
MaritalStatus unique values are : ['Single' 'Married' 'Divorced']
```

```
Over18 unique values are : ['Y']
```

```
OverTime unique values are : ['Yes' 'No']
```

## OBSERVATION:

- **Over18** has only one unique value ( Y ) which makes it irrelevant for our analysis

## DATA CLEANING

```
In [8]: # drop redundant columns
df.drop(['Over18', 'EmployeeCount', 'EmployeeNumber', 'StandardHours'],
        axis=1, inplace=True)
```

```
In [9]: df.shape
```

```
Out[9]: (1470, 31)
```

```
In [10]: # discretization of the Age column

bins = [17, 29, 39, 49, 60]
group_names = ['Youth', 'YoungAdult', 'MiddleAged', 'OldAdults']
df['AgeCategory'] = pd.cut(df['Age'], bins, labels=group_names)
```

```
In [11]: df['AgeCategory'].value_counts()
```

```
Out[11]: YoungAdult      622
MiddleAged      349
Youth          326
OldAdults       173
Name: AgeCategory, dtype: int64
```

```
In [12]: # Encodes the Attrition column
from sklearn.preprocessing import LabelEncoder

label_encoder = LabelEncoder()
df['AttritionEncoded'] = label_encoder.fit_transform(df['Attrition'])
```

## Exploratory Data Analysis(EDA)

```
In [13]: import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline
sns.set_style('darkgrid')
```

### 1. UNIVARIATE

### PIE CHART


```
In [14]: fig, axes = plt.subplots(2,3, figsize=(15, 8))
```

```

#Pie chart for Business Travel
df['BusinessTravel'].value_counts().plot(kind='pie', cmap='viridis', ax=axes[0][0])
#Pie chart for Attrition
df['Attrition'].value_counts().plot(kind='pie', ax=axes[0][1])
#Pie chart for Department
df['Department'].value_counts().plot(kind='pie', cmap='autumn', ax=axes[0][2])
#Pie chart for Gender
df['Gender'].value_counts().plot(kind='pie', cmap='coolwarm', ax=axes[1][0])
#Pie chart for Marital Status
df['MaritalStatus'].value_counts().plot(kind='pie', cmap='cividis', ax=axes[1][1])
#Pie chart for Overtime
df['OverTime'].value_counts().plot(kind='pie', cmap='inferno', ax=axes[1][2])

```

Out[14]: <AxesSubplot:ylabel='OverTime'>



## OBSERVATIONS:

- Most of employee rarely travel and most of them dont travel
- Most of the employee have No Attrition
- Most of the employee are from the Research and Development department followed by Sales then Human Resources
- There are more Male than Female
- Most of the employee are married
- Most of the employee dont work overtime

```



In [15]: fig, axes = plt.subplots(1,2, figsize=(15, 4))

sns.histplot(df[df['Attrition'] == 'No']['Age'], bins=30, kde=True, ax=axes[0])
sns.histplot(df[df['Attrition'] == 'Yes']['Age'], bins=30, kde=True, ax=axes[1], color='br')

axes[0].set_title('Distribution of Age of No Attrition')
axes[1].set_title('Distribution of Age of Yes Attrition')

```

Out[15]: Text(0.5, 1.0, 'Distribution of Age of Yes Attrition')





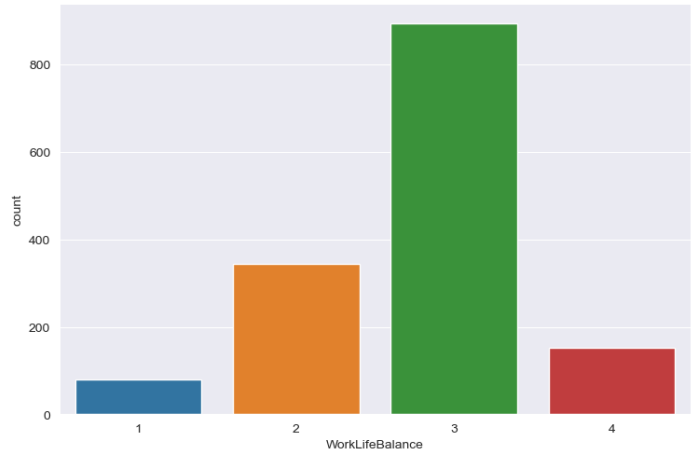



### OBSERVATION:

The Age Distribution looks the same for both category of Attrition

```
In [16]: fig, axs = plt.subplots(3,2, figsize=(15, 15))

sns.countplot(x=df['JobLevel'],hue=df['Attrition'],ax=axs[0][0])
sns.countplot(x=df['JobSatisfaction'],ax=axs[0][1])
sns.countplot(x=df['RelationshipSatisfaction'],hue=df['AgeCategory'], ax=axs[1][0])
sns.countplot(x=df['WorkLifeBalance'],ax=axs[1][1])
sns.countplot(x=df['StockOptionLevel'],hue=df['AgeCategory'],ax=axs[2][0])
sns.countplot(x=df['WorkLifeBalance'],ax=axs[2][1])

plt.tight_layout()
```

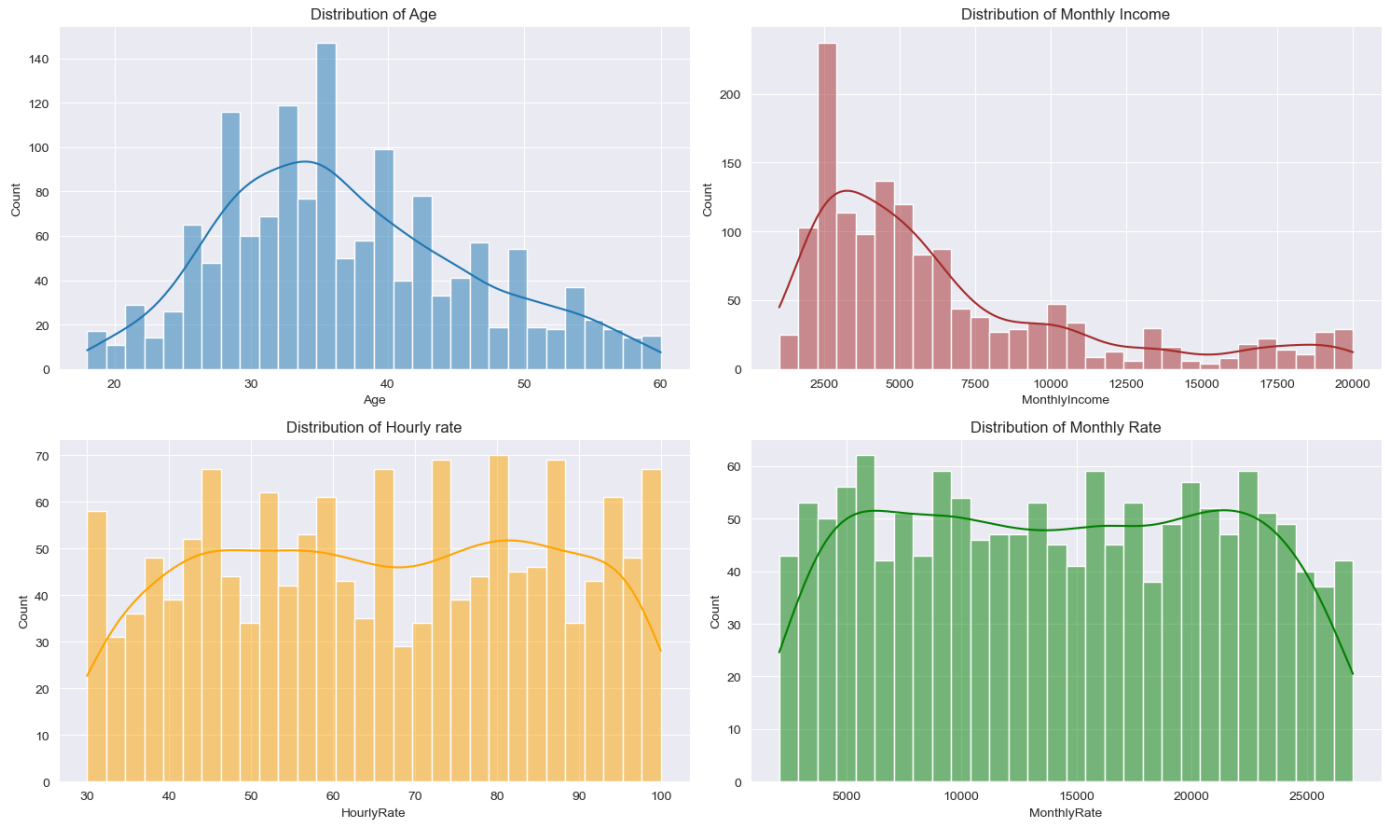


```
In [17]: fig, axs = plt.subplots(2,2, figsize=(15, 9))

sns.histplot(df['Age'], bins=30, kde=True, ax=axs[0][0])
sns.histplot(df['MonthlyIncome'], bins=30, kde=True, ax=axs[0][1], color='brown')
sns.histplot(df['HourlyRate'], bins=30, kde=True, ax=axs[1][0], color='orange')
sns.histplot(df['MonthlyRate'], bins=30, kde=True, ax=axs[1][1], color='green')

axs[0][0].set_title('Distribution of Age')
axs[0][1].set_title('Distribution of Monthly Income')
axs[1][0].set_title('Distribution of Hourly rate')
axs[1][1].set_title('Distribution of Monthly Rate')

plt.tight_layout()
```





## 2. Bivariate

```
In [18]: fig, axs = plt.subplots(2,2, figsize=(15, 9))

sns.countplot(x=df['BusinessTravel'],hue=df['Gender'], ax=axs[0][0])
sns.countplot(x=df['JobSatisfaction'], hue=df['Gender'], ax=axs[0][1])
sns.countplot(x=df['BusinessTravel'], hue=df['MaritalStatus'], ax=axs[1][0])
sns.countplot(x=df['JobSatisfaction'],hue=df['MaritalStatus'], ax=axs[1][1])

axs[0][0].set_title('Distribution of Business Travel')
axs[0][1].set_title('Distribution of Job Satisfaction')
```


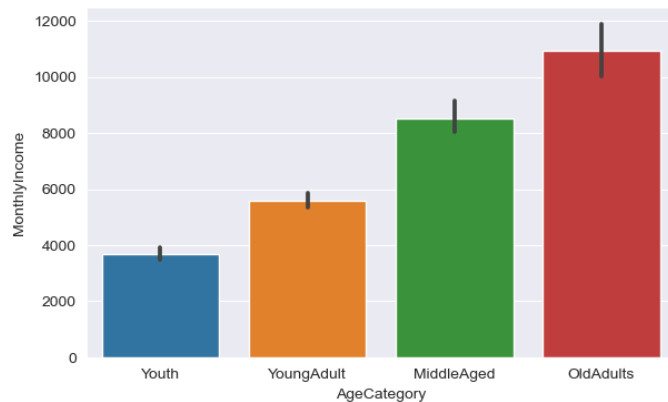


```
Out[18]: Text(0.5, 1.0, 'Distribution of Job Satisfaction')
```



```
In [19]: fig, axs = plt.subplots(2,2, figsize=(15, 9))

sns.barplot(x=df['Education'], y=df['MonthlyIncome'], estimator=np.mean, ax=axs[0][0])
sns.barplot(x=df['YearsWithCurrManager'], y=df['MonthlyIncome'], estimator=np.mean, ax=axs[0][1])
sns.barplot(x=df['AgeCategory'], y=df['MonthlyIncome'], estimator=np.mean, ax=axs[1][0])
sns.barplot(x=df['Department'], y=df['MonthlyIncome'], estimator=np.mean, ax=axs[1][1])

Out[19]: <AxesSubplot: xlabel='Department', ylabel='MonthlyIncome'>
```



**OBSERVATION:**




- The mean Monthly income increases with Education , Age and Years with current manager
- The sales department has the mean highest monthly income followed by Human resources then Research and development

```
In [20]: fig, axs = plt.subplots(2,1, figsize=(15, 9))

sns.barplot(x=df['YearsAtCompany'], y=df['MonthlyIncome'], estimator=np.mean, ax=axs[0])
sns.barplot(x=df['YearsInCurrentRole'], y=df['MonthlyIncome'], estimator=np.mean, ax=axs
```

```
Out[20]: <AxesSubplot:xlabel='YearsInCurrentRole', ylabel='MonthlyIncome'>
```



## OBSERVATION:

- The mean monthly income increases as the number of years in the company increases
- The mean monthly income increases as the number of years in the current role increases

## STATISTICAL ANALYSIS

```
In [21]: import scipy.stats as stats
```

```
In [22]: alpha = 0.05
```

## CHI-SQUARE TEST OF INDEPENDENCE

### 1. Attrition vs Gender

- H0: There is no association between Attrition and Gender
- H1: There is an association between Attrition and Gender

```
In [23]: # Create a contingency table of Attrition and Gender
crosstab = pd.crosstab(df['Attrition'], df['Gender'])
```

crosstab

Out[23]:    **Gender**   **Female**   **Male**

**Attrition**

**No**        501    732

**Yes**        87     150

In [24]:    *# Perform a chi-square test for association between Attrition and Gender*  
chi2, p\_value, \_, \_ = stats.chi2\_contingency(crosstab)

```
print("Chi-Square Test Results:")  
print("Chi-Square:", chi2)  
print("p-value:", p_value)
```

```
#reject H0?  
print('\nReject H0?: ')  
if p_value < alpha:  
    print('True')  
else:  
    print('False')
```

Chi-Square Test Results:  
Chi-Square: 1.1169671241970975  
p-value: 0.29057244902890855

Reject H0?:  
False

## CONCLUSION:

Since we do not reject the null hypothesis, the conclusion is that there is enough evidence to support the claim that Attrition does not depend on gender

## 2.Attrition vs Business Travel

- H0: There is no association between Attrition and Business Travel
- H1: There is an association between Attrition and Business Travel

In [25]:    *# Create a contingency table of Attrition and Business Travel*  
crosstab = pd.crosstab(df['Attrition'], df['BusinessTravel'])  
crosstab

Out[25]:    **BusinessTravel**   **Non-Travel**   **Travel\_Frequently**   **Travel\_Rarely**

**Attrition**

**No**            138                    208                    887

**Yes**            12                        69                    156

In [26]:    *# Perform a chi-square test for association between Attrition and Business Travel*  
chi2, p\_value, \_, \_ = stats.chi2\_contingency(crosstab)

```
print("Chi-Square Test Results:")  
print("Chi-Square:", chi2)  
print("p-value:", p_value)
```

```
#reject H0?
```

```

print('\nReject H0?: ')
if p_value < alpha:
    print('True')
else:
    print('False')

```

Chi-Square Test Results:  
Chi-Square: 24.182413685655174  
p-value: 5.608614476449931e-06

Reject H0?:  
True

## CONCLUSION:

The conclusion is that there is enough evidence to support the claim that Attrition is related to how frequent an employee travels

### 3. ATTRITION vs DEPARTMENT

- H0: There is no association between Attrition and Department
- H1: There is an association between Attrition and Department

```

In [27]: # Create a contingency table of Attrition and Department
crosstab = pd.crosstab(df['Attrition'], df['Department'])
crosstab

```

Out[27]: **Department**   **Human Resources**   **Research & Development**   **Sales**

Attrition			
No	51	828	354
Yes	12	133	92

```

In [28]: # Perform a chi-square test for association between Attrition and Department
chi2, p_value, _, _ = stats.chi2_contingency(crosstab)

```

```

print("Chi-Square Test Results:")
print("Chi-Square:", chi2)
print("p-value:", p_value)

```

```

#reject H0?
print('\nReject H0?: ')
if p_value < alpha:
    print('True')
else:
    print('False')

```

Chi-Square Test Results:  
Chi-Square: 10.79600732241067  
p-value: 0.004525606574479633

Reject H0?:  
True

## CONCLUSION:

The conclusion is that there is enough evidence to support the claim that Attrition is related to the department where the employee works

## ATTRITION VS MARITAL STATUS

- H0: There is no association between Attrition and Marital Status
- H1: There is an association between Attrition and Marital Status

```
In [29]: # Create a contingency table of Attrition and Marital Status
crosstab = pd.crosstab(df['Attrition'], df['MaritalStatus'])
crosstab
```

```
Out[29]: MaritalStatus  Divorced  Married  Single
```

Attrition			
No	294	589	350
Yes	33	84	120

```
In [30]: # Perform a chi-square test for association between Attrition and Marital Status
chi2, p_value, _, _ = stats.chi2_contingency(crosstab)
```

```
print("Chi-Square Test Results:")
print("Chi-Square:", chi2)
print("p-value:", p_value)
```

```
#reject H0?
print('\nReject H0?: ')
if p_value < alpha:
    print('True')
else:
    print('False')
```

```
Chi-Square Test Results:
Chi-Square: 46.163676540848705
p-value: 9.45551106034083e-11
```

```
Reject H0?:
True
```

## CONCLUSION:

The conclusion is that there is enough evidence to support the claim that Attrition is related to the Marital Status of the employee

## ATTRITION VS OVERTIME

- H0: There is no association between Attrition and Overtime
- H1: There is an association between Attrition and Overtime

```
In [31]: # Create a contingency table of Attrition and Overtime
crosstab = pd.crosstab(df['Attrition'], df['OverTime'])
crosstab
```

```
Out[31]: OverTime  No  Yes
```

Attrition		
No	944	289
Yes	110	127

```
In [32]: # Perform a chi-square test for association between Attrition and Overtime
chi2, p_value, _, _ = stats.chi2_contingency(crosstab)

print("Chi-Square Test Results:")
print("Chi-Square:", chi2)
print("p-value:", p_value)

#reject H0?
print('\nReject H0?: ')
if p_value < alpha:
    print('True')
else:
    print('False')
```

```
Chi-Square Test Results:
Chi-Square: 87.56429365828768
p-value: 8.15842372153832e-21
```

```
Reject H0?:
True
```

## CONCLUSION:

The conclusion is that there is enough evidence to support the claim that Attrition is related to whether the employee works overtime or not